



Training Deep Neural Networks



[YouTube Playlist](#)

Maziar Raissi


Assistant Professor

Department of Applied Mathematics

University of Colorado Boulder

maziar.raissi@colorado.edu

Tabular Data



Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.5	2.5	4.0	1.3	versicolor
5.0	2.0	3.5	1.0	versicolor
6.7	2.5	5.8	1.8	virginica
5.6	3.0	4.5	1.5	versicolor
5.2	2.7	3.9	1.4	versicolor
5.0	3.5	1.3	0.3	setosa
6.4	2.7	5.3	1.9	virginica
6.4	2.8	5.6	2.2	virginica
5.1	3.8	1.6	0.2	setosa
5.1	3.7	1.5	0.4	setosa

– Number of observations: 150 – Number of variables: 5

The goal of statistics is to come up with a “model” that “explains” the data.
The goal of machine (deep) learning is to come up with an “algorithm” that “performs” well on some “test” data.

$x_i \in \mathbb{R}^d, y_i \in \{1, 2, \dots, k\}, i = 1, \dots, N \rightarrow$ training data

$x'_i \in \mathbb{R}^d, y'_i \in \{1, 2, \dots, k\}, i = 1, \dots, N' \rightarrow$ validation data

$x''_i \in \mathbb{R}^d, y''_i \in \{1, 2, \dots, k\}, i = 1, \dots, N'' \rightarrow$ test data

Performance metric (e.g., accuracy)

$\hat{y} = f_{\alpha, \beta}(x) \rightarrow$ prediction of model $f_{\alpha, \beta}$ at input x

$\alpha \rightarrow$ hyper-parameters of the model

$\beta \rightarrow$ parameters of the model

$\hat{y}''_i = f_{\alpha^*, \beta^*}(x''_i) \rightarrow$ predictions of the “optimal” model on the test data

$$\frac{\sum_{i=1}^{N''} \mathbb{1}(\hat{y}''_i = y''_i)}{N''} \rightarrow \text{accuracy}$$

Model/Algorithm

$$f_{\alpha, \beta}(x) = \arg \min_{j=1, \dots, k} p_{\alpha, \beta}^{(j)}(x)$$

$p_{\alpha, \beta}^{(j)}(x) \rightarrow j$ -th element of the probability vector $p_{\alpha, \beta}(x) \in \mathbb{R}^k$

$$\sum_{j=1}^k p_{\alpha, \beta}^{(j)}(x) = 1 \text{ and } p_{\alpha, \beta}^{(j)}(x) \geq 0, \forall j = 1, \dots, k.$$

Training

$$\mathcal{L}_{\alpha}(\beta) = - \sum_{i=1}^N \log p_{\alpha, \beta}^{(y_i)}(x_i) \rightarrow \text{loss function (negative log likelihood)}$$

$$\beta^* = \arg \min_{\beta} \mathcal{L}_{\alpha}(\beta) \rightarrow \text{given } \alpha \text{ (i.e., } \beta^* \text{ is a function of } \alpha)$$

Validation

$\hat{y}'_i = f_{\alpha, \beta^*(\alpha)}(x'_i) \rightarrow$ predictions of the model on the validation data

$$\alpha^* = \arg \min_{\alpha} \frac{\sum_{i=1}^{N'} \mathbb{1}(\hat{y}'_i = y'_i)}{N'}$$

Multi-Layer Perceptron (MLP)

$$p_{\alpha, \beta}(x) = \text{softmax}(W \underbrace{\text{ReLU}(Vx + a)}_{=:h} + b)$$

An overview of gradient descent optimization algorithms


[YouTube Playlist](#)

```
#include
<stdio.h>
int main()
{
    printf("Hello")
;
    return 0;
}
```

Source code



Compiler

```
100010101010101
000100101010111
011111100110000
001011001101010
010111011100011
011111001111000
000110011110101
010010010101000
```

Executable code

Deep Learning: An algorithm that writes an algorithm

Source Code: Data (examples/experiences)

Compiler: Deep Learning

Executable Code: Deployable Model

Deep: Function Compositions $f_L \circ f_{L-1} \circ \dots \circ f_2 \circ f_1$

Learning: Loss Function, Back-propagation, and Gradient Descent

$$\min_{\theta} L(\theta)$$

$L(\theta) \approx J(\theta) \rightarrow$ noisy estimate of the objective function (e.g., due to mini-batching)

Stochastic Gradient Descent

$$\theta_{t+1} = \theta_t - \gamma \nabla_{\theta} J(\theta_t)$$

$$J : \mathbb{R}^d \rightarrow \mathbb{R}$$

one back-propagation (fast)

$$\nabla_{\theta} J : \mathbb{R}^d \rightarrow \mathbb{R}^d$$

$$\nabla_{\theta} J : \mathbb{R}^d \rightarrow \mathbb{R}^d$$

$$\nabla_{\theta}^2 J : \mathbb{R}^d \rightarrow \mathbb{R}^{d \times d}$$

d back-propagations (slow)

Momentum

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta_t)$$

$$\theta_{t+1} = \theta_t - v_t$$

Nesterov Accelerated Gradient (NAG)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta_t - \gamma v_{t-1})$$

Adagrad

$$g_t = \nabla_{\theta} J(\theta_t)$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t$$

$$G_t = \sum_{\tau=1}^t g_{\tau} \odot g_{\tau}$$

Adadelta

$$E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma) g_t^2$$

$$\theta_{t+1} = \theta_t + \Delta \theta_t$$

$$\Delta \theta_t = - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

units don't match

$$\Delta \theta_t = - \frac{RMS[\Delta \theta]_{t-1}}{RMS[g]_t} g_t$$

RMSprop

$$E[g^2]_t = 0.9 E[g^2]_{t-1} + 0.1 g_t^2$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

Adam

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\left. \begin{aligned} \hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \\ \hat{v}_t &= \frac{v_t}{1 - \beta_2^t} \end{aligned} \right\} \text{bias-corrected}$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$

AdaMax

$$u_t = \max(\beta_2 u_{t-1}, |g_t|)$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{u_t} \hat{m}_t$$

Nadam

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} (\beta_1 \hat{m}_t + \frac{(1 - \beta_1) g_t}{1 - \beta_1^t})$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} (\beta_1 \hat{m}_{t-1} + \frac{(1 - \beta_1) g_t}{1 - \beta_1^t}) \rightarrow \text{Adam}$$



Questions?



[YouTube Playlist](#)
